

# Programming with Persistent Memory from Java

April 24, 2020

# Persistent Memory for Java

## Intel open-source libraries

- Low-Level Persistence Library (LLPL) -- version 1.0
- Java bindings to PMDK pmemkv library -- version 0.9
- Persistent Collections for Java (PCJ) -- experimental

## OpenJDK enhancements

- JEP 352 Persistent MappedByteBuffer -- in JDK14
- JEP 370 Foreign-Memory Access API -- incubator in JDK14

Links to these are on last slide

# Comparison of Libraries

	LLPL	Java bindings to pmemkv	PCJ	Mapped ByteBuffer	Memory Access API
Status	release 1.0	release 0.9	experimental	JDK14	incubator JDK14
Compatibility	JDK 8+	JDK 8+	JDK 8+	JDK14+	JDK14
Persistent data	heaps of memory blocks	key-value store	Java collections and other classes	ByteBuffers	structs, unions, arrays, etc.
Memory mgmt.	manual	manual	automatic	manual	manual
Thread-safe	heap: yes blocks: no	yes - optional	yes	no	no
Data integrity / consistency	developer- defined	developer- defined	ACID objects	developer- defined	developer- defined
Transactions	yes	yes - on puts	yes	no	no

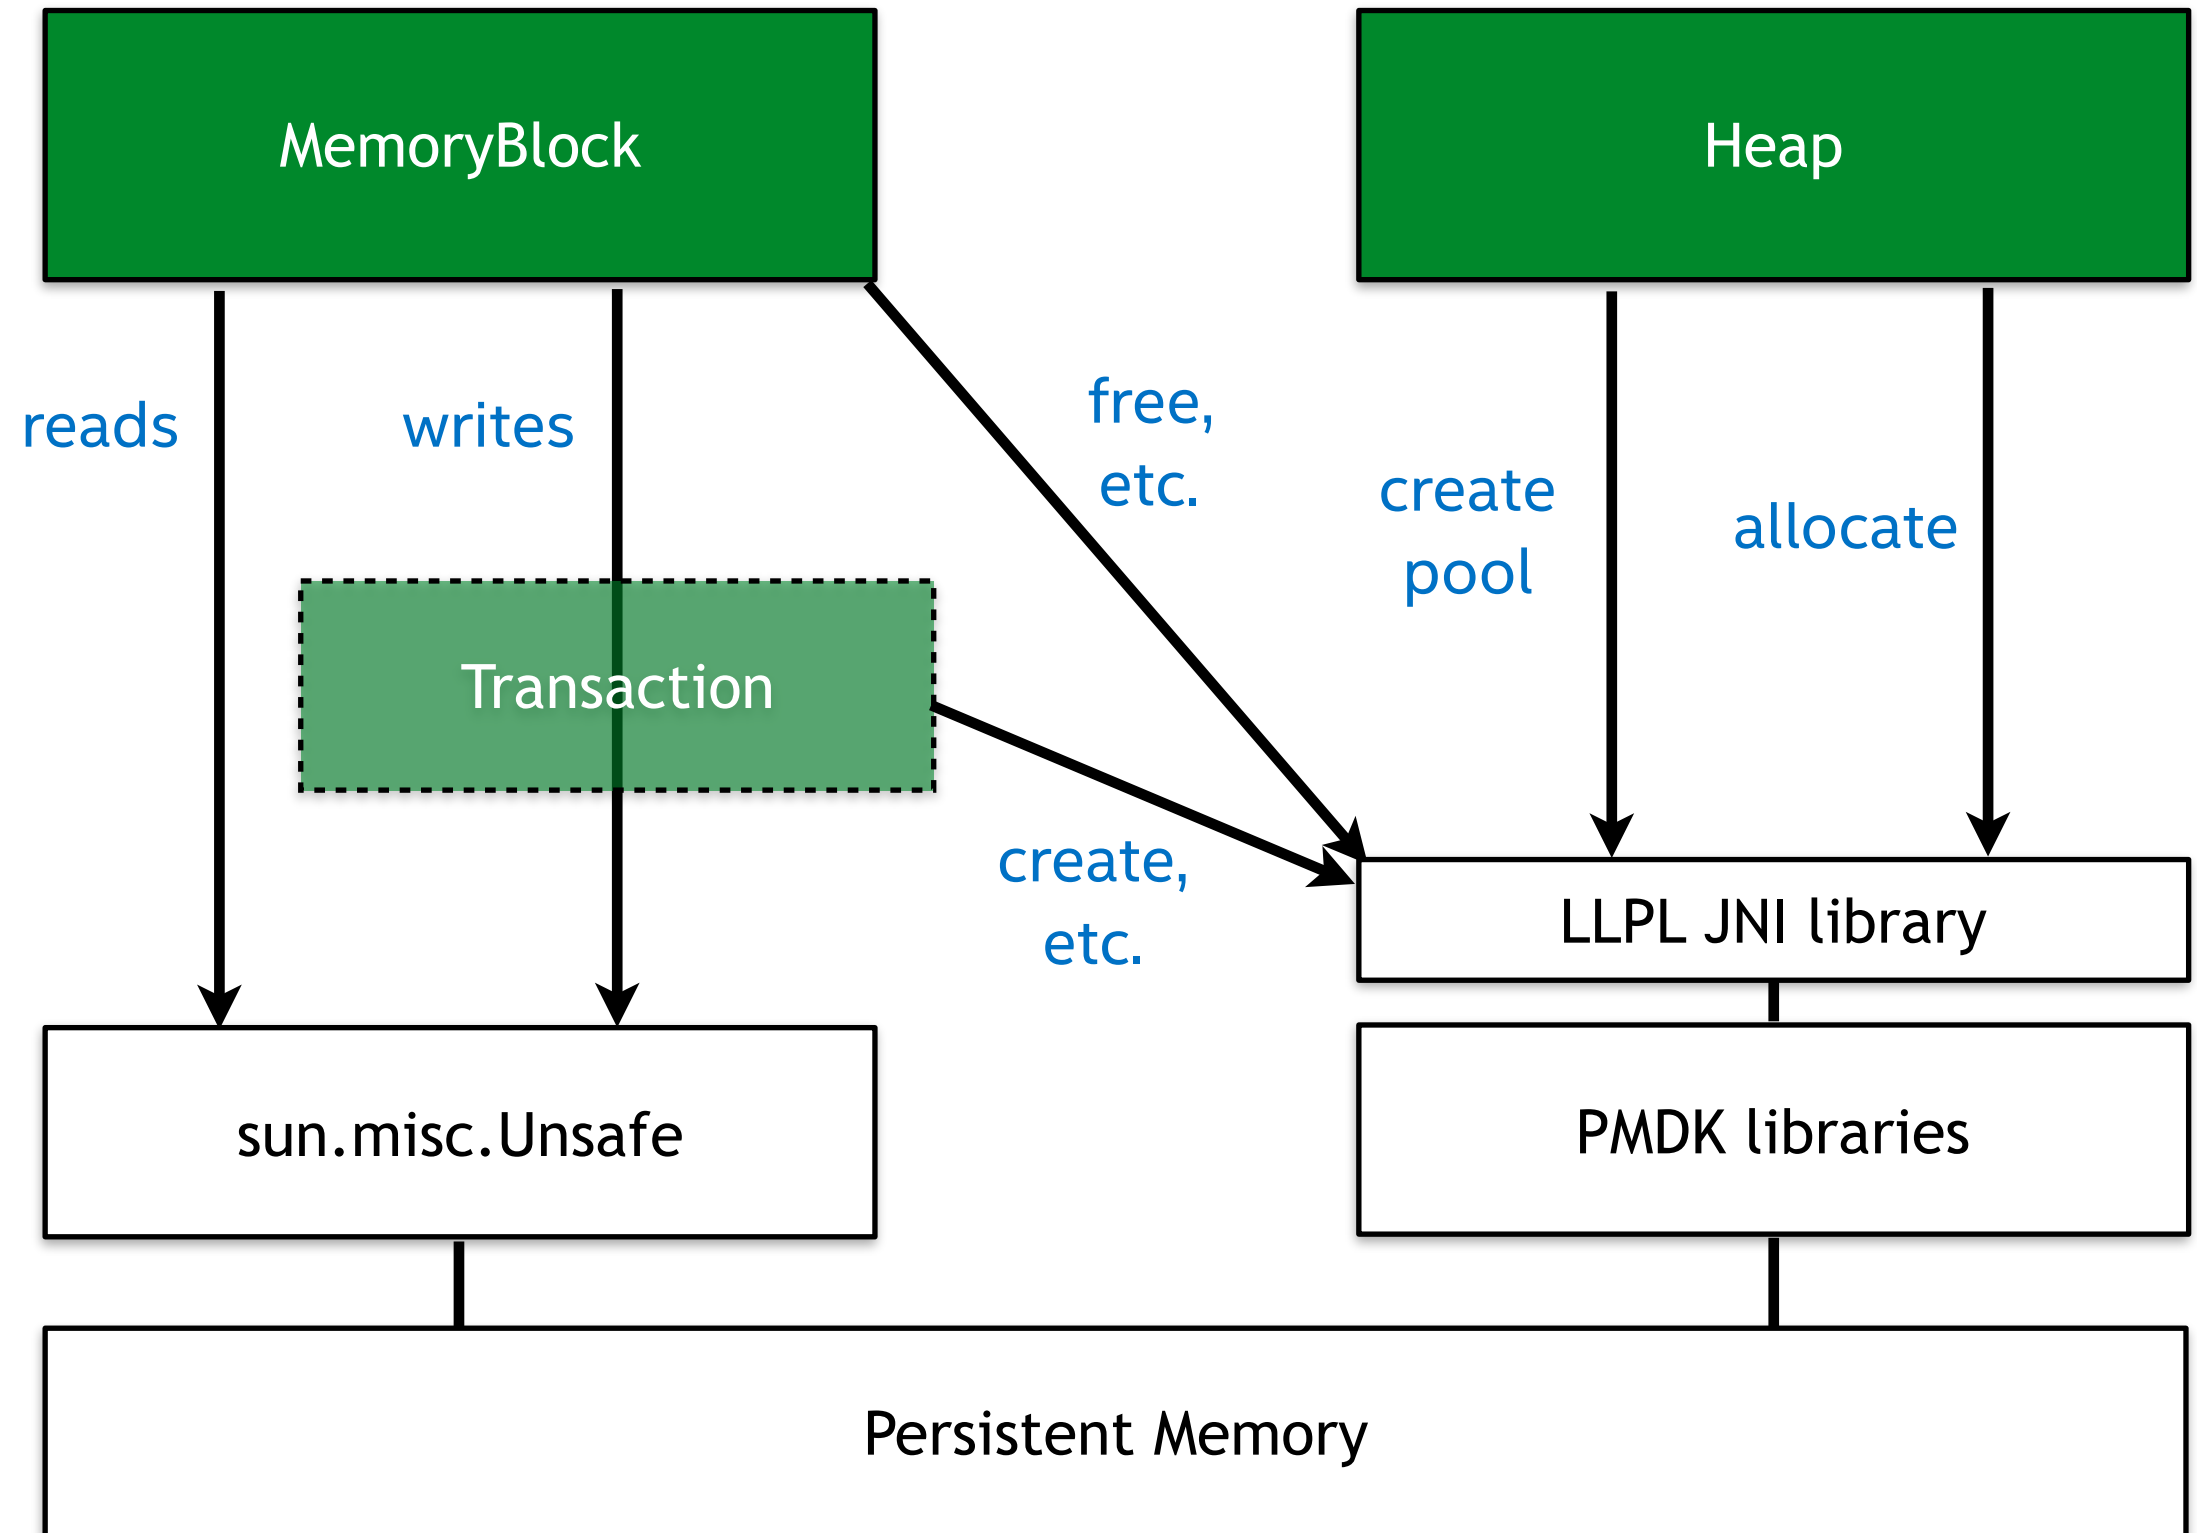
# LLPL Overview

- Intel open-source Java library for persistent memory programming
- Compatible with JDK 8+
- Version 1.0 released at end of 2019
- A component of the Persistent Memory Development Kit
- Depends on PMDK libraries (libpmem, libpmemobj)
- Supports building with Maven or Make
- MavenCentral availability in-progress
- Goals:
  - Flexible, high-performance access to pmem from Java
  - Easy to use, idiomatic Java API (despite low-level nature)
  - Suitable for direct use or as base for higher-level abstractions

# LLPL -- Three Primary Elements

- **Heap** -- a pool memory and an allocator for it
  - can create multiple heaps of almost any size
  - re-open a heap to access after restart
  - manual memory management
  - thread-safe API
- **MemoryBlock** -- an accessor for a block of allocated memory
  - low-level setters and getters indexed by offsets within a block
  - can refer to and recover a block long-term using stable (long) handles
  - write handles in another block to link blocks and data structures
  - non-thread-safe API
- **Transaction** -- can group writes for fail-safe data consistency
  - transaction state is thread-local
  - API integrates with try-catch and uses lambdas as bodies
  - nested transactions "flattened" to commit or abort together

# LLPL Implementation



# MemoryBlock API -- Reading and Writing

## Write methods:

1. setByte
2. setShort
3. setInt
4. setLong
5. setMemory
6. copyFromArray
7. copyFromMemoryBlock

## Read methods:

1. getByte
2. getShort
3. getInt
4. getLong
5. copyToArray

## Other methods:

- flush
- free
- addToTransaction

# Data Integrity and Consistency

- Being able to say something clear about the usability of heap data after an event such as:
  - a controlled exit
  - an unhandled exception
  - a power failure
- Durable changes to pmem leave heap data in a known-usable state if writes aren't interrupted, e.g. normal execution, controlled exit
- Transactional changes to pmem leave heap data in a known-usable state after any of the above events
- LLPL offers tools to create your own simple or not-so-simple policies



# Three Writes, Two Errors, Three Heaps

Three kinds of writes:

1. volatile: write data
2. durable: write data, **flush data from CPU cache to media**
3. transactional: **add data range to transaction** (back up data), write data

Two new kinds of programming errors:

1. durable write: forget to flush data from cache
2. transactional write: forget to add range to transaction before writing

Three kinds of heaps:

1. Heap: most flexible but both errors are possible
2. PersistentHeap: all durable, opt. transactional -- if code compiles error #1 is not present
3. TransactionalHeap: all transactional -- if code compiles neither error is present

# Heap Fragmentation

- Allocated handles are stable for the lifetime of the allocation
- Some allocate / free patterns can fragment the heap
- Same is true for off-heap DRAM; exacerbated here by long lived heaps
- Mitigations and solutions
  - reuse: choose sizes that enable reuse, e.g. when updating existing data
  - sub-allocation: application-managed sub-allocation of large blocks
  - compaction API (exploring now): developer requests that the heap compact a selection of handles returning new handles
  - storage service (like pmemkv): data is copied in and copied out of a data structure; service owns the memory and is free to move it to defragment

# LLPL Performance and Usability Optimizations

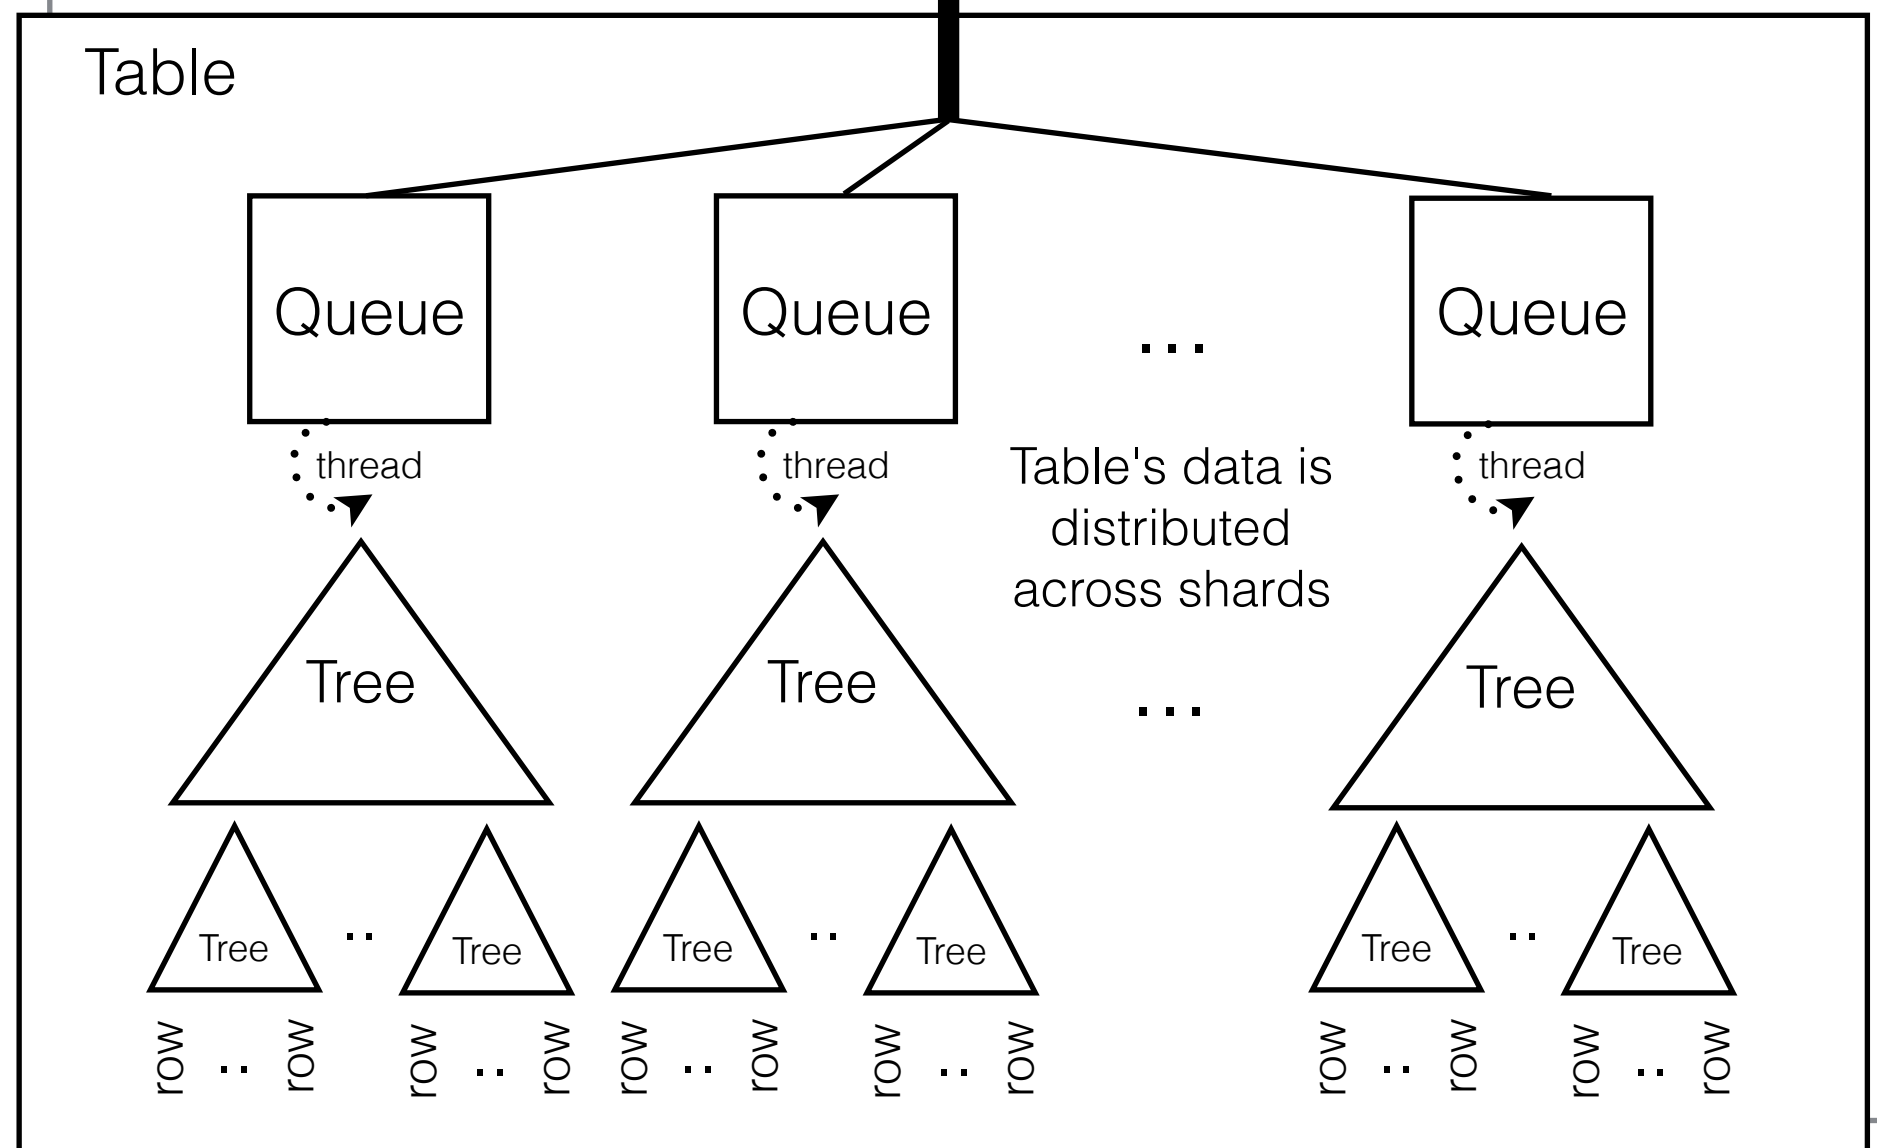
- All reads and writes done using Unsafe
- Minimized number of JNI calls
- Easy to use "ranged write" methods automate and minimize flush and transaction operations:
  - on initialization of new memory block
  - developer-defined range of offsets
- Automatically benefit from core PMDK optimizations, e.g.
  - allocator
  - transactions

# Example Use: Cassandra Pmem Storage Engine


<https://github.com/intel/cassandra-pmem>

All persistent data structures implemented using LLPL

## Storage Engine



# Cassandra Throughput



- using 4x P4510 NVMe SSDs
- using persistent memory storage engine

# LLPL Workshop and Repository Examples

## Workshop

1. getting started
2. sizing heaps
3. using other heaps
4. more on transactions
5. performance idioms
6. wrapping memory blocks

## Repository

1. int array
2. handle array
3. linked list
4. adaptive radix tree

# Going Forward

- Future of these open-source libraries is steered by customer feedback and requirements
- Examples of LLPL features / optimizations being considered:
  - construction-free allocation
  - re-positionable accessor
  - small library of data structures
  - handle tracking for liveness, etc.
- Please don't hesitate to contact us!

# Links

- Intel® Optane™ DC persistent memory
  - <https://www.intel.com/content/www/us/en/products/memory-storage/optane-dc-persistent-memory.html>
- Low Level Persistence Library (LLPL)
  - <https://github.com/pmem/llpl>
- Java bindings to PMDK pmemkv library -- version 0.9
  - <https://github.com/pmem/pmemkv-java>
- Persistent Collections for Java (PCJ) [Experimental]
  - <https://github.com/pmem/pcj>
- Persistent Memory Development Kit (PMDK)
  - <https://github.com/pmem/pmdk>
- JEP 370 - java.foreign Memory Access API
  - <https://openjdk.java.net/jeps/370>
- JEP 352 -- Non-Volatile Mapped Byte Buffers
  - <https://openjdk.java.net/jeps/352>
- JEP 316 -- Allocation of Java Heap on Alt. Memory Devices
  - <https://openjdk.java.net/jeps/316>
- Cassandra persistent memory storage engine
  - <https://github.com/intel/cassandra-pmem>